

# A Survey on Conversational AI in Task-oriented and Question-Answering tasks

Sanchit Nevgi

College of Information and Computer Sciences  
University of Massachusetts, Amherst

## Abstract

Ubiquitous deep learning approaches have led to the advent of neural Conversation AI systems, also known as dialogue systems. Modern dialogue system can be distinctly categorized into 3 modules — A Natural Language Understanding (NLU), Dialogue Manager (DM), and Natural Language Generation (NLG). In this survey, we investigate the evolution of dialogue systems, from hand-crafted rules and statistical approaches to recent end-to-end data-driven neural methods. Specifically we focus on task-oriented dialog systems, where the dialog agent works toward a specific goal, for example, booking movie tickets as well as open-domain question answering.

## 1 Introduction

Advancements in deep learning have led to improved performance in dialogue systems. A *dialogue agent* is an entity that you can interface with via voice/text that has certain capabilities. These capabilities can be broadly classified into three categories [Gao et al., 2019] as,

**Question & Answering:** QA agents allow users to query large-scale Knowledge Bases (KB-QA) or document collection in natural language (text-QA). In the real-world, text-QA agents are most commonly used in search engines such as Google, Bing, while KB-QA are widely used in voice assistants.

**Task completion:** Task-oriented systems work towards a well-specified goal, such as movie ticket booking, usually in a multi-turn fashion. The dialogue system keeps track of the dialogue state, uses information supplied by user to constrain search, prompts user for required information for task completion.

**Social dialogue:** The agent needs to converse seamlessly with users, provide useful recommendations. A good dialogue agent should demonstrate emotional connect with the user. For example, showing excitement when user shares good news, empathy when user is feeling sad, etc.

In this review, we study the evolution of the approaches used in building an effective dialogue agent, with an emphasis on task-oriented goals. Traditional approaches for building task-oriented dialogs relied on hand-crafting features [Core and Allen, 1997, Jelinek, 1976] and casting dialogue as an optimal decision making problem [Young et al., 2013, Kaelbling et al., 1998]. These approaches are unable to adapt to new domains, aren't robust to *user paraphrasing*, and semantic word similarity. Further, in these approaches, the system is comprised of several modules (NLU, NLG, etc) that are optimized independently shown in Figure 3. Recent approaches utilize neural network models that are trained end-to-end.

## 2 Dialogue as an Optimal Decision making task

We briefly review of Reinforcement Learning (RL) and Markov Decision Processes. A comprehensive overview in RL is out of scope for this review. Interested readers are referred to the excellent work by [Sutton and Barto, 1988, Kaelbling et al., 1996]

**Reinforcement Learning:** Reinforcement Learning (RL) is a learning paradigm where given an agent in an environment, the agent has to learn to take optimal decisions. Every action is followed by an immediate reward, which is a signal to the agent of its performance. Each iteration of the agent in the environment is described by,

1. The agent observes the environment's current state  $s_t$  and takes an action  $a_t$

2. The environment transitions to the next state  $s_{t+1}$ , distributed according to transition probability  $P(\cdot|s_t, a_t)$
3. The agent receives an immediate reward  $r_t$  post transition.

The agent-environment interaction is often modeled as a discrete-time Markov decision process, described next.

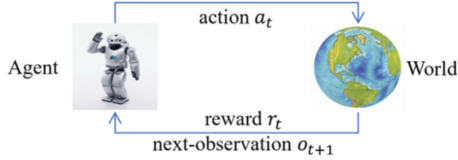


Figure 1: Reward system in RL environment

**Markov Decision Processes (MDP):** A Markov Decision Process is used to describe an environment for reinforcement learning, where the environment is fully observable. An MDP is a Markov Reward Process with decisions, it's an environment in which all states are Markov [Sutton et al., 1999].

**Partially Observable MDP (POMDP):** A Spoken Dialog System (SDS) can be cast as a Partially Observable MDP [Williams and Young, 2007, Young et al., 2013]. A POMDP can be developed to encompass a complete dialog system and serves as a basis for optimization. It can integrate uncertainty in the form of statistical distributions.

Formally, a POMDP is defined as a tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \gamma, b_0)$  where

- $\mathcal{S}$  is a set of all states
- $\mathcal{A}$  is a set of actions
- $\mathcal{T}$  defines a transition probability  $P(s_t|s_{t-1}, a_{t-1})$
- $\mathcal{R}$  defines the expected immediate real-valued reward,
- $\mathcal{O}$  is a set of observatoins
- $\mathcal{Z}$  defines an observation probability  $P(o_t|s_t, a_{t-1})$
- $\gamma$  is the geometric discount factor  $0 \leq \gamma \leq 1$
- $b_0$  is the initial belief state.

At each time step  $t$ , the world is in some unobserved state  $s_t$ . A distribution over possible states called belief states  $b_t$  is maintained where  $b_t(s_t)$  is the probability of state  $s_t$ . Based on this, the model




Figure 2: A Partially Observed Markov Decision Process

selects an action  $a_t$ , receives an observation  $o_{t+1}$ . Given an existing belief state  $b_t$ , the last system action  $a_t$ , and the new observation  $o_{t+1}$ , the new update belief state  $b_{t+1}$  is given by

$$b_{t+1}(s_{t+1}) = \eta P(o_{t+1}|s_{t+1}, a_t) \cdot \sum_{s_t} P(s_{t+1}|s_t, a_t) b_t(s_t) \quad (1)$$

where  $\eta$  is the normalization constant. The system action is determined by policy  $\pi$ . It is commonly represented by a deterministic mapping from belief states to actions  $\pi(b) \in \mathcal{A}$  or stochastically via a distribution over actions  $\pi(a|b) \in [0, 1]$

The discounted sum of rewards expected by starting in belief state  $b_t$ , using policy  $\pi$  is given by the *value function* as

$$V^\pi(b_t) = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots] \quad (2)$$

A similar quantity is the Q-function  $Q^\pi(b, a)$ , which gives the expected discounted sum of rewards if action  $a$  is taken given belief state  $b$ , followed by policy  $\pi$ . The *optimal policy*  $\pi^*$  is one that maximizes  $V^\pi$  to yield  $V^*$ , also known as the *Bellman optimality* equation. Including an explicit Bayesian model of uncertainty and by optimizing the policy via a reward-driven process, POMDPs provide a framework for modeling dialog managers [Young et al., 2010]. These models are optimized by maximizing the expected cumulative sum of rewards at each turn.

For real-world applications, the computation over all possible states becomes intractable and applying a direct solution is very difficult. Various approximation techniques have been proposed.

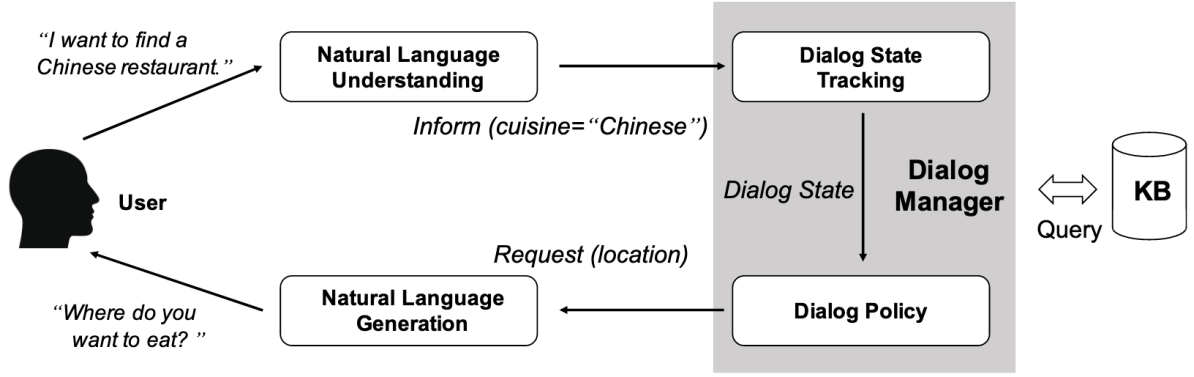


Figure 3: Modules in an NLP system

The principal components of a conversational Spoken Dialog System are shown in Figure. At each turn  $t$ , an SLU component converts each spoken input into an abstract semantic representation called *user dialog act*  $u_t$ . The *system* updates its internal state  $s_t$  and determines the next system act via a decision rule  $a_t = \pi(s_t)$ , known as *policy*. The system act is converted to speech using NLG component.

The posterior probability of the belief state after each user input is updated via Bayesian inference in a process known as *belief monitoring*. The design of the belief state allows user behaviour to be captured as model priors and the inference process is able to exploit the full distribution of recognition hypothesis. By maintaining a belief distribution over all states, the system is effectively pursuing all possible dialog paths in parallel, choosing the next action based on all the most likely state but all all states. The explicit representation of state and policy action allow dialog design to be incorporated by associating rewards with state-action pairs.

**Belief State Representation and Monitoring:** In a practical application, the state must encode the user’s utterance  $u_t$ , the history context  $h_t$  and the user’s goal  $g_t$ . Therefore, the state is a combination of three components  $s_t = f(g_t, u_t, h_t)$ . Reasonable independence assumptions are made regarding the conditional independencies of the model.

**Challenges:** The *state-action* space is extremely large, which makes real-time Bayesian inference rather challenging. The exact policy learning for POMDPs is intractable and approximation techniques have to be used. Further, learning requires real-world users which are expensive to acquire and time-consuming. Hence, building capa-

ble user simulators are required for learning.

### 3 Natural Language Generation

Natural Language Generation is vital to convert a communication goal, selected by the Dialogue Manager, into natural language. Traditionally rule/template-based methods were used for generation, with the transition to data-driven approaches. We briefly review recent approaches to NLG.

#### 3.1 SEQ2SEQ models

Also known as encoder-decoder architecture, these models are fundamental to language modeling tasks, such as machine translation and dialogue response generation. They are typically implemented using recurrent architectures such as RNNs and their counterpart, LSTMs [Hochreiter and Schmidhuber, 1997]. In these models, semantic representations of the input text tokens are passed through the *encoder*, through a series of transforms, yielding a hidden representation of the input. The output is generated token-by-token using this representation, which is known as the *decoder*.

#### 3.2 Contextualized word embeddings and Transformers

Recurrent models faced the issue of keeping track of long-term context, which prompted the use of attention. [Vaswani et al., 2017] showed that only using the attention mechanism was sufficient for model generation performance, called as the Transformer architecture. Subsequently, the BERT model [Devlin et al., 2019] used bi-directional transformers to get pre-trained contextualized word embeddings that possessed deep language understanding.

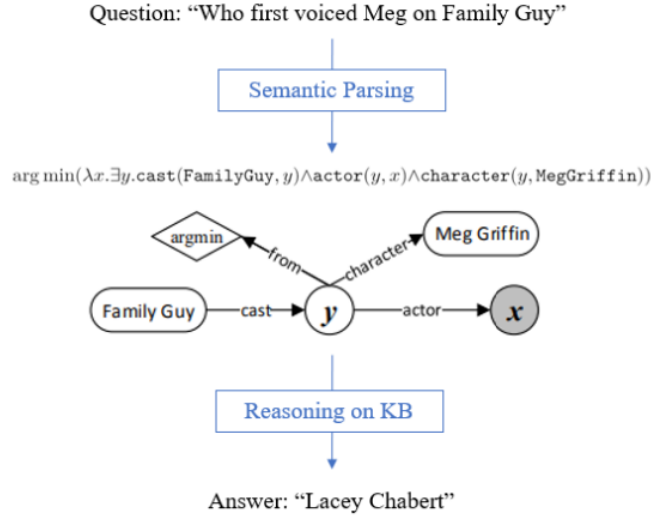
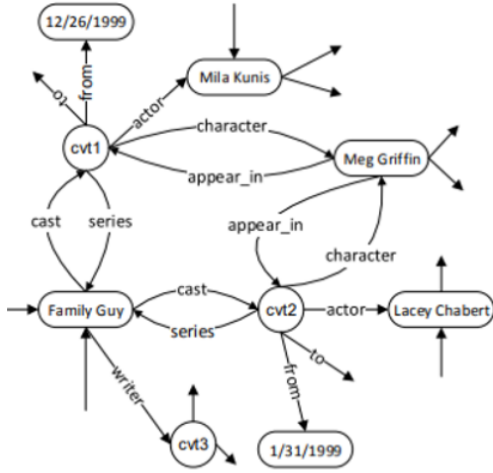


Figure 4: (Left) Semantic Parsing for KB-QA on a subgraph of Freebase. (Right) A question in its logical form using  $\lambda$ -calculus

## 4 Question Answering

Traditionally Conversational Question Answering relied on creating and maintaining hand-engineered rules. With the rise of deep learning, methods have been developed to automate this process. There are two main types of QA agents - KB-QA agents query a large-scale Knowledge Base (KB) to retrieve information. text-QA agents use Information Retrieval techniques on a document collection, to obtain a ranked list of relevant documents. Google Search, Bing are examples of text-QA agents.

### 4.1 Knowledge Bases

A Knowledge Base is a structured store of facts, usually in graph format. Typically, a KB consists of *subject-predicate-object* triples  $(s, r, t)$ , where  $s, t \in \mathcal{E}$  are entities and  $r \in \mathcal{R}$  is the relation between the entities. Some popular KBs are Freebase [Bollacker et al., 2008], which has since been integrated into Google Search, ConcpetNet [Speer et al., 2016] which is huge graph consisting of approximately 8m concepts and relations between them.

State-of-the-art symbolic approaches to reasoning over KB are based on *semantic-parsing* [Berant et al., 2013, Einolghozati et al., 2019], where a question is mapped to its formal meaning representation and then translated to a KB query. Applying symbolic KB-QA to a very large KB is challenging due to — Paraphrasing in Natural Language and

the search complexity.

## 5 Task-Oriented Systems

With recent advancements in on-device machine learning, voice assistants such as Google Assistant, Siri, Alexa have promulgated. These systems assist users in completing tasks such as making restaurant reservations, booking movie tickets, setting an alarm. These tasks differ from the QA in that they work toward a specific goal over multiple turns. We review the various components in Task-oriented systems with emphasis on recent neural approaches.

**Slot-filling dialogues:** The user and the system converse with the goal of *slot-filling*. The system must collect all the necessary information from the user to formulate an appropriate query. In the movie ticket booking domain, some examples of slots are *movie-name*, *time*, *price*, *location*. Slots are *informable* if the user can provide the *slot-value*, used to constrain the search; or slots are *requestable*, where the user can ask the value of the slot from the system (eg *phone\_number*).

**Dialogue Act:** The interaction between the user and the system, known as the dialogue act, can be modeled as a agent-environment problem in Reinforcement Learning, where the user is equivalent to the environment and the system is the agent. Briefly, at each turn, the agent keeps track of the dialogue state, then takes an action based on the



current state. The user responds with an utterance and an immediate reward is computed to measure the quality of the turn. We elaborate these steps in more detail further.

The architecture of task-oriented dialog systems can be broadly divided into two categories — Pipeline and End-to-end approaches. The pipeline system consists of 4 components,

1. Natural Language Understanding
2. Dialogue State Tracking
3. Dialogue Policy
4. Natural Language Generation

NLU maps to structured semantic representation. A popular schema for semantic representation is *dialog act*, which consists of intent and slot-values. The NLU task is mapped into intent-detection and slot extraction. Intent detection is framed as classification problem and slot-extraction is framed as sequence labeling task.




Figure 5: POMDP-based dialog manager

## 5.1 Learning End-to-End Task-oriented Dialog

Learning task-oriented dialogs end-to-end requires defining a user simulator and a dialogue agent. Next, we define baseline tasks to measure neural methods’s performances. Traditional dialog systems in task-oriented dialogs require lot of domain-specific handcrafting, which makes it difficult to scale. These limitations are mitigated in end-to-end trained systems. [Bordes and Weston, 2016] shows that, compared to hand-crafted slot-filling baseline, end-to-end dialog system based on *Memory networks* [Sukhbaatar et al., 2015] can reach promising, but imperfect performance and can even perform non-trivial operations.

Successful goal-oriented dialog systems model coversion as partially observable Markov deicision processes. Require hand-crafted features for state and action space representations, restricting to narrow domains. To measure the performance

of task-oriented dialog systems, 5 common tasks have been proposed. They are,

- Task 1 - Issuing API calls
- Task 2 - Updating API calls
- Task 3 - Displaying options
- Task 4 - Providing extra information
- Task 5 - Full dialogs

In [Bordes and Weston, 2016], the authors use a restaurant KB, which has name, cuisine, location, price range, address, phone number. Each example in the dataset comprises of dialog utterances from user and bot, as well as API calls and the resulting facts. They split the KB into two, disjoint sets of *cuisines*, and *locations*, to test the models capability to handle Out-of-Vocabulary scenario. The dialog systems are evaluated in a ranking, not a generation setting. At each turn of dialog, they test whether they can predict bot utterances by selecting candidates.

Further, the model is evaluated on real concierge service where the dialogs are shorter, and the vocabulary more diverse. Some of the findings are that the set of user requests is much wider compared to dataset, from managing restaurant reservation to asking for recommendations. Also, the users do not stay focused on the request. The facts about restaurants are not structured like KB and finally the users and operators make typos, spelling and grammatical errors.

Models:

Rule-based systems - Used hand-crafted rules such as word matches, positions in dialog, entity detections, dialog state.

Classic IR models - TF-IDF match - For each possible candidate response, compute the matching score between input and response. Score is TF-IDF weighted cosine similarity between the bag-of-words of input and response. Nearest neighbour - Find most similar conversation in the training set and use that response. Word overlap as scoring method. Sorted by decreasing co-occurrence frequency.

Supervised embedding models - Predict next response given the previous conversation. Scored for input  $x$  as  $f(x, y) = (Ax)^T By$ . Trained with margin ranking loss

Memory Networks -

Word embeddings fail with entities. Embeddings use approx word match. Cannot handle OOV

words. Solution - match type features. Augment vocab with 7 words (cuisine, location, etc)

## 5.2 Building End-To-End dialogue systems using Generative Hierarchical Neural Network Models

[Serban et al., 2015] investigate the task of open-domain, conversational dialogue systems. They make Use heirarchical recurrent encoder-decoder neural network. POMDPs are popular, and demonstratively work well for limited examples. Further, non-task-oriented sytems can be used to develop user simulators, which can be further used to generate dialogue datasets to train POMDP models. The dialogue response generation can be modeled as a translation task. Although it is more difficult due to numerous plausible responses and lack of phrase alignment between user utterance and response.

Dialog is viewed as sequence of utterances where each utterance comprises of tokens. The architecture consitsits of an ENCODER-RNN, which maps each utterance to an utterance vector. It is the hidden state obtained after last token is processed. A Context-RNN iteratively processes each utterance vector. The hidden state of the CONTEXT-RNN represents state of the dialogue system. It is passed to the DECODER-RNN to predict the next utterance. The HRED is hypothesized to be superior to RNN because (1) the Context-RNN allows model to represent common ground between speakers and (2) Number of computational steps between utterances is reduced. [Serban et al., 2015] show that Hierarchical Recurrent model outperforms n-gram based model and the baseline neural networks. They infer that using a large external monologue corpus to initialize word embeddings, and a large non-dialogue corpus for pre-training improves model performance. Finally they conclude that MAP outputs produce generic, but conversationally acceptable responses while stochastic samples from model produced diverse dialogues.

## 5.3 Neural Belief Tracker: Data-Driven Dialogue State Tracking

[Mrksic et al., 2016] proposes an end-to-end trained DST module known at the *Neural Belief Tracker*. A *belief tracker*, estimates the user's goal at every step of the dialog. They mitigate two issues with traditional approaches to building DST, namely that NLU models require large amount of training data and hand-crafting lexicons for captur-

ing liguistic variation in users's language is cumbersome.

Using recent advances in *representation learning*, Neural Belief Tracking (NBT) framework reasons over pre-trained word vectors, learning to compose them into distributed representations of user utterances and dialogue context. The *dialogue state tracking* component serves to interpret user input and update the *belief state*, which is the system's internal representation of teh state of the conversation. The dialogue system is supported by *domain ontology*, which describes the range of user intents the system can process. The ontology defines a collection of slots and the values each slot can take.

The task is non-trivial due to lexical variation, dynamics of context and noisy Automated Speech Recognition (ASR) output. Traditional approaches use separate modules to handle lexical variability in single dialogue turn. However the turn-level SLU and cross-turn DST can be coalesced into a single model, but they rely on manually constructed semantic dictionaries.

Some systems use template-based matching systems while others train independent binary models that decide if slot-value pair was expressed in user utterance. SLU has been treated as sequence labeling problem.

The proposed Neural Belief Tracking model uses pre-trained vectors. The input consists of the system dialogue acts preceding the user input, the user utterance, a single slot-value pair it needs to make a decision about. To perform belief tracking, NBT model iterates over all candidate slot-value pairs and decides which ones have just been expressed by the user. "*I'm looking for good pizza*" entails *FOOD=ITALIAN*.

**Representation Learning module:** The vector embeddings for unigram, bigram and trigram are concatenated and projected to a fixed-size representation. The canditate slot-value pair and user utterance interact through the *semantic decoding* module. The slot and value representation are concatenated, projected to the same size at user utterance vector and finally a similarity score is computed.

To conclude, the NBT couples Spoken Language Understanding and Dialogue State Tracking without relying on hand-crafted semantic lexicons. Further, the model performance improves with the semantic quality of underlying word vectors.

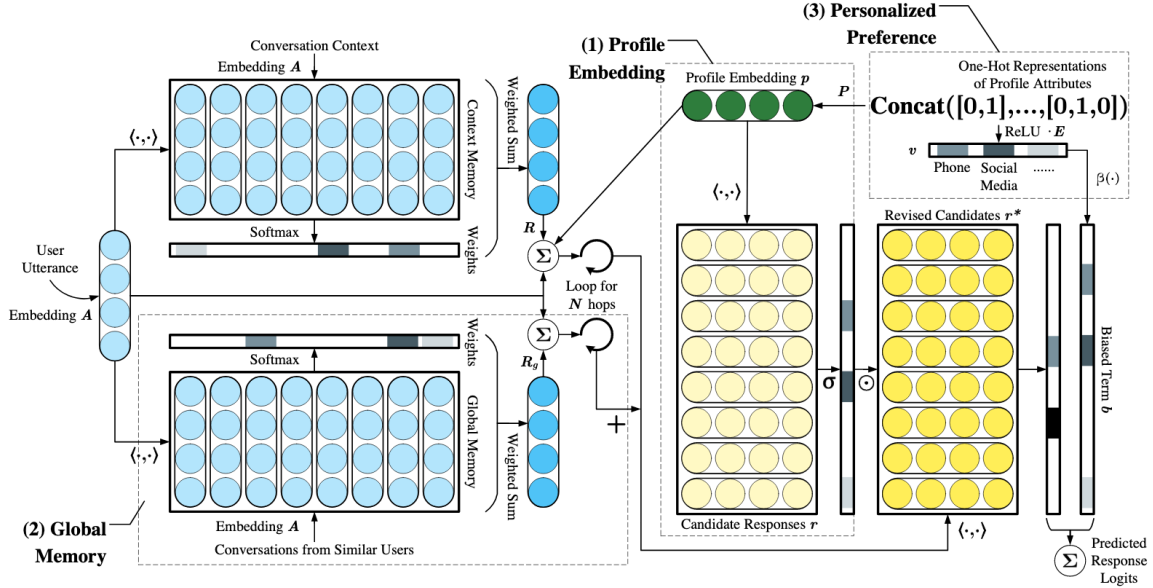


Figure 6: PERSONALIZED MEMN2N architecture

## 6 Learning Personalized End-to-End Goal-Oriented Dialog

Existing works on dialog systems only take into account the conversation content, neglecting the user personality. [Luo et al., 2019] proposes using a PROFILEMODEL, which encodes user profiles into distributed embeddings and a PREFERENCEMODEL captures user preferences over KB.

Modern approaches train dialogue systems end-to-end [Vinyals and Le, 2015, Sukhbaatar et al., 2015]. They are directly trained on past dialogs, without domain assumptions. Some common limitations are the model’s (1) inability to adjust language style according to the user (2) lack of dynamic conversation policy based on interlocutor’s profile (3) inability to handle ambiguities in user requests. Psychology studies have shown that during a dialog, humans adapt dialog style according to the interlocutor, which enhances conversational efficiency. An existing challenge is how to integrate user profile to generate personalized messages.

The PROFILE MODEL learns user personalities with distributed profile representations and uses a global memory to store conversation context from users with similar profiles. The PREFERENCE MODEL learns user preferences among ambiguous candidates by building a connection between user profile and knowledge base. These models are integrated into MEMN2N network. A common approach to leveraging personality in re-

cent works is using a conditional language model as the response decoder. Existing literature [Li et al., 2016] focuses on personality in social bots with limited work in task-oriented personality dialogues.

**Notation:** A user profile is represented by  $n$  attributes  $\{(k_i, v_i)\}_{i=1}^n$ , where  $k_i$  and  $v_i$  denote the key and value of the  $i$ -th attribute. For example  $\{(\text{Gender}, \text{Male}), (\text{Age}, \text{Young})\}$ . The values are represented as one-hot vectors. A complete *user profile*  $\hat{a}$  is obtained by the concatenation of all such one-hot representations.

**ProfileModel:** A profile embedding of the *user profile* is obtained by linear transformation  $p = P\hat{a}$ . The query  $q$  in a MEMN2N model plays a key role in reading the memory and choosing the response. At each turn, the profile embedding is incorporated into the query update as

$$q_{i+1} = q_i + o_i + p \quad (3)$$

where  $o_i$  is the output at the  $i$ -th hop.

**Global Memory:** Users with similar profile should expect a similar response for a particular request. The personalized information of similar users is incorporated into global memory module. The global memory is similar in structure to the MEMN2N model, with the difference that the contents of memory are history utterances from similar users instead of current conversation.

## 7 Evaluation

Existing measures of evaluation either prevent reproducibility, as different human annotators (Mechanical Turks) would have different experiences or that the metrics don't correlate with human judgements. Automatic evaluation of the dialogue systems is still an active area of research.

Individual components of the system, such as the Natural Language Understanding module, may be evaluated using traditional techniques such as accuracy, precision/recall, BLEU scores. However, to get a holistic view of the performance of the entire system, they are insufficient. A common metric used is the *task success rate*, which tracks whether the dialog agent correctly fulfilled the task (booked movie tickets, reserved a table, etc). [Williams and Young, 2007].

Human evaluation is the most accurate measure of the system performance. However, it is expensive to operate on a large scale. Recently, [Adiwardana et al., 2020] proposed a new human-evaluation metric known as Sensibility and Specificity Average (SSA). The authors used a 2.6B parameter SEQ2SEQ model with Evolved Transformer, trained on 40B words. The SSA metric was shown to highly correlate to *perplexity*, which is easier to compute.

## 8 Challenges in Neural Task-oriented dialog

Neural approaches for task completion is advancing at a rapid pace [Zhang et al., 2020]. The frontier in model performance is pushed every day by bigger and better models. However, it is important to consider the data efficiency of these models to facilitate dialog system modeling in low-resource settings. A major feature of task-oriented dialog is its emphasis on multi-turn state-action dynamics. Faster and exact policy learning for modeling multi-turn dynamics is essential. Similarly, building models that adapt to new domains with reasonable accuracy are crucial. Lastly, integrating domain ontology knowledge into the dialog model remains a challenge.

## References

Jianfeng Gao, Michel Galley, and Lihong Li. Neural approaches to conversational ai. *Found. Trends Inf. Retr.*, 13:127–298, 2019.

- Mark G. Core and James F. Allen. Coding dialogs with the damsl annotation scheme. 1997.
- Frederick Jelinek. Speech recognition by statistical methods. 1976.
- Steve J. Young, Milica Gasic, Blaise Thomson, and Jason D. Williams. Pomdp-based statistical spoken dialog systems: A review. *Proceedings of the IEEE*, 101:1160–1179, 2013.
- Leslie Pack Kaelbling, Michael L. Littman, and Anthony R. Cassandra. Planning and acting in partially observable stochastic domains. *Artif. Intell.*, 101:99–134, 1998.
- Richard S. Sutton and Andrew G. Barto. Reinforcement learning: An introduction. *IEEE Transactions on Neural Networks*, 16:285–286, 1988.
- Leslie Pack Kaelbling, Michael L. Littman, and Andrew W. Moore. Reinforcement learning: A survey. *ArXiv*, cs.AI/9605103, 1996.
- Richard S. Sutton, Doina Precup, and Satinder P. Singh. Between mdps and semi-mdps: A framework for temporal abstraction in reinforcement learning. *Artif. Intell.*, 112:181–211, 1999.
- Jason D. Williams and Steve J. Young. Partially observable markov decision processes for spoken dialog systems. *Comput. Speech Lang.*, 21:393–422, 2007.
- Steve J. Young, Milica Gasic, Simon Keizer, François Mairesse, Jost Schatzmann, Blaise Thomson, and Kai Yu. The hidden information state model: A practical framework for pomdp-based spoken dialogue management. *Comput. Speech Lang.*, 24:150–174, 2010.
- Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Computation*, 9:1735–1780, 1997.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *NIPS*, 2017.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *NAACL-HLT*, 2019.
- Kurt D. Bollacker, C. J. Evans, Praveen Paritosh, Tim Sturge, and Jamie Taylor. Freebase: a collaboratively created graph database for structuring human knowledge. In *SIGMOD Conference*, 2008.
- Robyn Speer, Joshua Chin, and Catherine Havasi. Conceptnet 5.5: An open multilingual graph of general knowledge. In *AAAI*, 2016.
- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on freebase from question-answer pairs. In *EMNLP*, 2013.



800	Arash Einolghozati, Panupong Pasupat, Sonal Gupta,	850
801	Rushin Shah, Mrinal Mohit, Mike Lewis, and Luke	851
802	Zettlemoyer. Improving semantic parsing for task	852
803	oriented dialog. <i>ArXiv</i> , abs/1902.06000, 2019.	853
804	Antoine Bordes and Jason Weston. Learning end-to-	854
805	end goal-oriented dialog. <i>ArXiv</i> , abs/1605.07683,	855
806	2016.	856
807	Sainbayar Sukhbaatar, Arthur Szlam, Jason Weston,	857
808	and Rob Fergus. End-to-end memory networks. In	858
809	<i>NIPS</i> , 2015.	859
810	Iulian Serban, Alessandro Sordoni, Yoshua Bengio,	860
811	Aaron C. Courville, and Joelle Pineau. Building end-	861
812	to-end dialogue systems using generative hierarchi-	862
813	cal neural network models. In <i>AAAI</i> , 2015.	863
814	Nikola Mrksic, Diarmuid Ó Séaghdha, Tsung-Hsien	864
815	Wen, Blaise Thomson, and Steve J. Young. Neu-	865
816	ral belief tracker: Data-driven dialogue state track-	866
817	ing. <i>ArXiv</i> , abs/1606.03777, 2016.	867
818	Liangchen Luo, Wenhao Huang, Qi Zeng, Zaiqing Nie,	868
819	and Xu Sun. Learning personalized end-to-end goal-	869
820	oriented dialog. In <i>AAAI</i> , 2019.	870
821	Oriol Vinyals and Quoc V. Le. A neural conversational	871
822	model. <i>ArXiv</i> , abs/1506.05869, 2015.	872
823	Jiwei Li, Michel Galley, Chris Brockett, Georgios P.	873
824	Spithourakis, Jianfeng Gao, and William B. Dolan.	874
825	A persona-based neural conversation model. <i>ArXiv</i> ,	875
826	abs/1603.06155, 2016.	876
827	Daniel De Freitas Adiwardana, Minh-Thang Luong,	877
828	David R. So, Jamie Hall, Noah Fiedel, Romal	878
829	Thoppilan, Zi Yang, Apoorv Kulshreshtha, Gau-	879
830	rav Nemade, Yifeng Lu, and Quoc V. Le. To-	880
831	wards a human-like open-domain chatbot. <i>ArXiv</i> ,	881
832	abs/2001.09977, 2020.	882
833	Zheng Zhang, Ryuichi Takanobu, Minlie Huang, and	883
834	Xiaoyan Zhu. Recent advances and challenges in	884
835	task-oriented dialog system, 2020.	885
836	Adarsh Kumar, Peter Ku, Anuj Kumar Goyal, Angeliki	886
837	Metallinou, and Dilek Hakkani-Tur. Ma-dst: Multi-	887
838	attention based scalable dialog state tracking, 2020.	888
839		889
840		890
841		891
842		892
843		893
844		894
845		895
846		896
847		897
848		898
849		899